

Combining in Spark:

1. Pseudo-code Spark Scala - Follower Count:

```
rddGroupByKey(sc: SparkContext, input: String, output: String):  
  inputRDD = sc.textFile(input)  
  pairsRDD = inputRDD.map(row => (row.split(",")(1), 1))  
  groupedRDD = pairsRDD.groupByKey()  
  followersCountRDD = groupedRDD.map(x => (x._1, x._2.sum))  
  followersCountRDD.saveAsTextFile(output)
```

```
rddReduceByKey(sc: SparkContext, input: String, output: String):  
  inputRDD = sc.textFile(input)  
  pairsRDD = inputRDD.map(row => (row, 1))  
  reducedRDD = pairsRDD.reduceByKey(_+_)  
  reducedRDD.saveAsTextFile(output)
```

```
rddFoldByKey(sc: SparkContext, input: String, output: String):  
  inputRDD = sc.textFile(input)  
  pairsRDD = inputRDD.map(row => (row.split(",")(1), 1))  
  foldRDD = pairsRDD.foldByKey(0)(_+_)  
  foldRDD.saveAsTextFile(output)
```

```
rddAggregateByKey(sc: SparkContext, input: String, output: String):  
  inputRDD = sc.textFile(input)  
  pairsRDD = inputRDD.map(row => (row.split(",")(1), 1))  
  aggregatorRDD = pairsRDD.aggregateByKey(0)(_+_,_+_)  
  aggregatorRDD.saveAsTextFile(output)
```

```
groupByDataset(spark: SparkSession, input: String, output: String):  
  dataSchema = StructType(Array(  
    StructField("follower", LongType, nullable = false),  
    StructField("user", LongType, nullable = false)))  
  inputDF = spark.read.format("csv").schema(dataSchema).load(input)  
  groupDF = inputDF.groupBy("user").count()  
  groupDF.write.csv(output)
```

2. *Reporting toDebugString() for RDD-based programs*

- *RDD-G*

```
20/02/28 21:06:38 INFO SparkContext: RDD's recursive dependencies:
(3) MapPartitionsRDD[5] at saveAsTextFile at SparkCombiner.scala:16 []
  | MapPartitionsRDD[4] at map at SparkCombiner.scala:15 []
  | ShuffledRDD[3] at groupByKey at SparkCombiner.scala:13 []
+- (3) MapPartitionsRDD[2] at map at SparkCombiner.scala:12 []
    | input MapPartitionsRDD[1] at textFile at SparkCombiner.scala:11 []
    | input HadoopRDD[0] at textFile at SparkCombiner.scala:11 []
```

- *RDD-R*

```
20/02/28 21:10:49 INFO SparkContext: RDD's recursive dependencies:
(3) MapPartitionsRDD[4] at saveAsTextFile at SparkCombiner.scala:24 []
  | ShuffledRDD[3] at reduceByKey at SparkCombiner.scala:22 []
+- (3) MapPartitionsRDD[2] at map at SparkCombiner.scala:21 []
    | input MapPartitionsRDD[1] at textFile at SparkCombiner.scala:20 []
    | input HadoopRDD[0] at textFile at SparkCombiner.scala:20 []
```

- *RDD-F*

```
20/02/28 21:16:01 INFO SparkContext: RDD's recursive dependencies:
(3) MapPartitionsRDD[4] at saveAsTextFile at SparkCombiner.scala:32 []
  | ShuffledRDD[3] at foldByKey at SparkCombiner.scala:30 []
+- (3) MapPartitionsRDD[2] at map at SparkCombiner.scala:29 []
    | input MapPartitionsRDD[1] at textFile at SparkCombiner.scala:28 []
    | input HadoopRDD[0] at textFile at SparkCombiner.scala:28 []
```

- *RDD-A*

```
20/02/28 21:23:01 INFO SparkContext: RDD's recursive dependencies:
(3) MapPartitionsRDD[4] at saveAsTextFile at SparkCombiner.scala:40 []
  | ShuffledRDD[3] at aggregateByKey at SparkCombiner.scala:38 []
+- (3) MapPartitionsRDD[2] at map at SparkCombiner.scala:37 []
    | input MapPartitionsRDD[1] at textFile at SparkCombiner.scala:36 []
    | input HadoopRDD[0] at textFile at SparkCombiner.scala:36 []
```

3. Reporting explain() for Dataset-based programs:

- *DSET*

```

== Parsed Logical Plan ==
Aggregate [user#1L], [user#1L, count(1) AS count#7L]
+- Relation[follower#0L,user#1L] csv

== Analyzed Logical Plan ==
user: bigint, count: bigint
Aggregate [user#1L], [user#1L, count(1) AS count#7L]
+- Relation[follower#0L,user#1L] csv

== Optimized Logical Plan ==
Aggregate [user#1L], [user#1L, count(1) AS count#7L]
+- Project [user#1L]
   +- Relation[follower#0L,user#1L] csv

== Physical Plan ==
*(2) HashAggregate(keys=[user#1L], functions=[count(1)], output=[user#1L, count#7L])
+- Exchange hashpartitioning(user#1L, 200)
   +- *(1) HashAggregate(keys=[user#1L], functions=[partial_count(1)], output=[user#1L, count#12L])
      +- *(1) FileScan csv [user#1L] Batched: false, Format: CSV, Location: InMemoryFileIndex[file:/Users/akshitjain,
], PushedFilters: [], ReadSchema: struct<user:bigint>

```

4. State clearly which of the programs performs aggregation before shuffling, and which does not.
- RDD-R, RDD-F and RDD-A perform aggregation before shuffling, whereas RDD-G, and DSET don't.

Join Implementation:

RS-R (Spark Scala Pseudocode):

```
MAX_FILTER = 50000

rddReduceSideJoin(sc: SparkContext, input: String, output: String):
  rowRDD = sc.textFile(input)
  filterRDD = rowRDD.filter(row => {
    rowVals = row.split(",")
    from = rowVals(0)
    to = rowVals(1)
    from < MAX_FILTER && to < MAX_FILTER
  })
  fromRDD = filterRDD.map(row => {
    rowVals = row.split(",")
    from = rowVals(0)
    to = rowVals(1)
    (from, to)
  })
  toRDD = filterRDD.map(row => {
    rowVals = row.split(",")
    from = rowVals(0)
    to = rowVals(1)
    (to, from)
  })
  counter = Accumulator()
  path2RDD = joinOnKey(fromRDD, toRDD).map(_._2)
  triangleRDD = joinOnKey(path2RDD, fromRDD).map(_._2)
  triangleRDD.foreach { x => if (x._1 == x._2) counter.add(1) }
  print(counter.value / 3)
}

joinOnKey(fromRDD, toRDD):
  joinedRDD = fromRDD.join(toRDD)
  joinedRDD
}
```

RS-D (Spark Scala Pseudocode):

```
MAX_FILTER = 50000;

dsReduceSideJoin(spark: SparkSession, input: String, output: String)
  customSchema = StructType(Array(
    StructField("follower_id", LongType, nullable = false),
    StructField("followee_id", LongType, nullable = false)
  ))
  followerDS = spark.read.format("csv").schema(customSchema)
    .load(input)
    .where(follower_id < MAX_FILTER && followee_id < MAX_FILTER)

  path2DS = followerDS.as("a").joinWith(followerDS.as("b"), a.followee_id === b.follower_id)

  triangleDS = path2DS.as("a").joinWith(followerDS.as("b"),
    a._1.follower_id === b.followee_id && a._2.followee_id === b.follower_id)
  print(triangleDS.count() / 3)
```

Rep-R (Spark Scala Pseudocode):

```
MAX_FILTER = 50000;

rddReplicatedJoin(sc: SparkContext, input: String, output: String):
  rowRDD = sc.textFile(input)
  filterRDD = rowRDD.filter(row => {
    rowValues = row.split(",")
    from = rowValues(0)
    to = rowValues(1)
    from.toLong < MAX_FILTER && to.toLong < MAX_FILTER
  })
  fromRDD = filterRDD.map(row => {
    rowValues = row.split(",")
    from = rowValues(0)
    to = rowValues(1)
    (from, to)
  })
  toRDD = filterRDD.map(row => {
    rowValues = row.split(",")
    from = rowValues(0)
    to = rowValues(1)
    (from, to)
  })
  counter = Accumulator()
  edgesRDD = fromRDD.collect().groupBy { case (from, to) => to }
  broadcastVal = sc.broadcast(edgesRDD)
  path2 = toRDD.mapPartitions(iter => {
    iter map {case (from, to) =>
      if (broadcastVal.value.get(from).isDefined) {
        arr = broadcastVal.value(from)
        arr.foreach { case (mapFrom, mapTo) =>
          if (broadcastVal.value.get(mapFrom).isDefined) {
            arr1 = broadcastVal.value(mapFrom)
            arr1.foreach { case (f, t) =>
              if (to.equals(f)) {
                counter.add(1)
              }
            }
          }
        }
      }
    }
  })
  print(counter.value / 3)
```

Rep-D (Spark Scala Pseudocode):

```

MAX_FILTER = 50000;

dsReplicatedJoin(spark: SparkSession, input: String, output: String):
  customSchema = StructType(Array(
    StructField("follower_id", LongType, nullable = false),
    StructField("followee_id", LongType, nullable = false)))

  followersDS = spark.read.format("csv").schema(customSchema).
    load(input)
    .where(follower_id < MAX_FILTER && followee_id < MAX_FILTER)

  path2DS =
    followersDS.as("a").join(broadcast(followersDS).as("b"),
      a.followee_id === b.follower_id).select(a.follower_id, b.followee_id)

  triangleDS =
    path2DS.as("a").joinWith(followersDS.as("b"),
      a.followee_id === b.followee_id && a.followee_id === b.follower_id)

  print(triangleDS.count() / 3)

```

Comparing Running Time:

Configuration	Small Cluster Result	Large Cluster Result
RS-R, MAX = 15000	Running time: 17 mins Triangle count: 1096146	Running time: 14 mins, Triangle count: 1096146
RS-D, MAX = 60000	Running time: 19 mins, Triangle count: 2.3801662E7	Running time: 12 mins, Triangle count: 2.3801662E7
Rep-R, MAX = 10000	Running time: 21 mins, Triangle count: 520296	Running time: 18 mins, Triangle count: 520296
Rep-D, MAX = 180000	Running time: 10 mins, Triangle count: 7.3323616E7	Running time: 8 mins, Triangle count: 7.3323616E7

Links to the Log files:

/Users/akshitjain/Desktop/Northeastern-University/Large-Scale/spark-twitter-data/logs/RS-R-Small.txt
/Users/akshitjain/Desktop/Northeastern-University/Large-Scale/spark-twitter-data/logs/RS-R-Large.txt

/Users/akshitjain/Desktop/Northeastern-University/Large-Scale/spark-twitter-data/logs/RS-D-Small.log
/Users/akshitjain/Desktop/Northeastern-University/Large-Scale/spark-twitter-data/logs/RS-D-Large.log

/Users/akshitjain/Desktop/Northeastern-University/Large-Scale/spark-twitter-data/logs/Rep-R-Small.log
/Users/akshitjain/Desktop/Northeastern-University/Large-Scale/spark-twitter-data/logs/Rep-R-Large.log

/Users/akshitjain/Desktop/Northeastern-University/Large-Scale/spark-twitter-data/logs/Rep-D-Small.log
/Users/akshitjain/Desktop/Northeastern-University/Large-Scale/spark-twitter-data/logs/Rep-D-Large.log

Links to the Output files:

/Users/akshitjain/Desktop/Northeastern-University/Large-Scale/spark-twitter-data/output/RS-R-Small-Output.txt
/Users/akshitjain/Desktop/Northeastern-University/Large-Scale/spark-twitter-data/output/RS-R-Large-Output.txt

/Users/akshitjain/Desktop/Northeastern-University/Large-Scale/spark-twitter-data/output/RS-D-Small-Output.txt
/Users/akshitjain/Desktop/Northeastern-University/Large-Scale/spark-twitter-data/output/RS-D-Large-Output.txt

/Users/akshitjain/Desktop/Northeastern-University/Large-Scale/spark-twitter-data/output/Rep-R-Small-Output.txt
/Users/akshitjain/Desktop/Northeastern-University/Large-Scale/spark-twitter-data/output/Rep-R-Large-Output.txt

/Users/akshitjain/Desktop/Northeastern-University/Large-Scale/spark-twitter-data/output/Rep-D-Small-Output.txt
/Users/akshitjain/Desktop/Northeastern-University/Large-Scale/spark-twitter-data/output/Rep-D-Large-Output.txt